



Generative Artificial Intelligence

Parameter Efficient Fine-Tuning



Outline

1. PEFT Introduction
2. PEFT Theory
 - Intrinsic Dimensionality
3. PEFT Current Development & Method
 - Adapters
 - Prompt Tuning
 - Bitfit
 - LoRA
 - MAM Adapters
 - S4




PEFT

Introduction



LLM Full Finetune Predicament

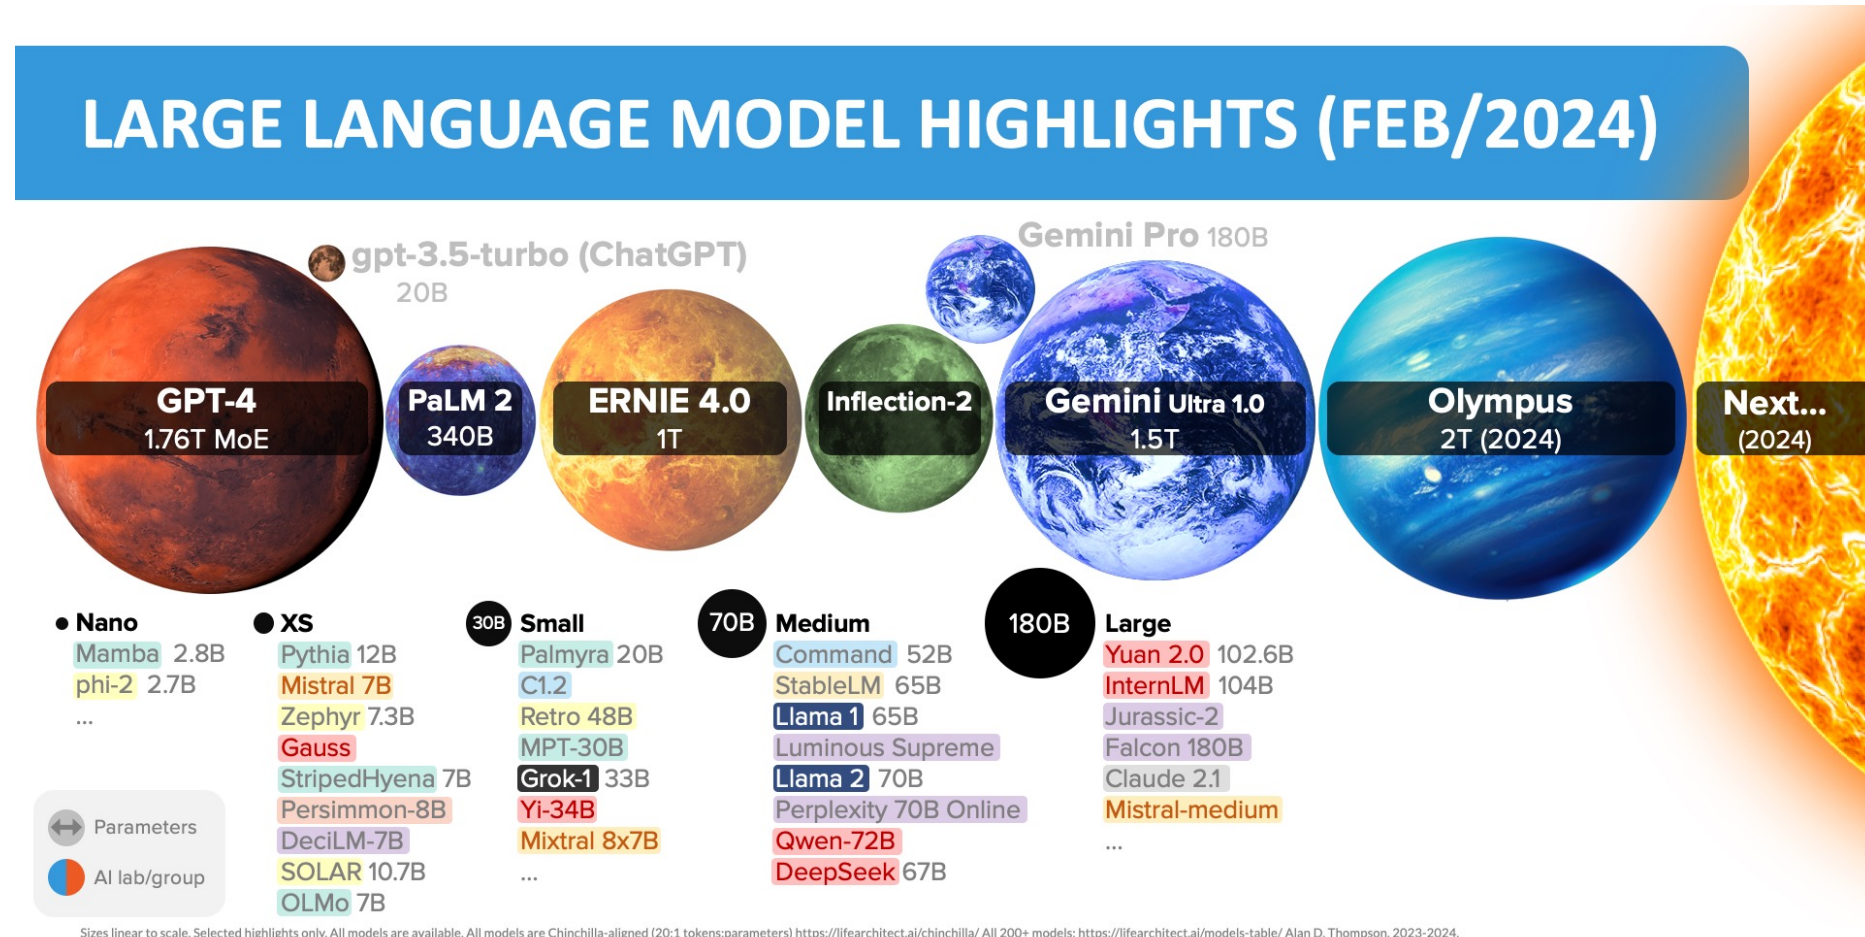


Model Name	Model Size	Developer
PaLM	540 Billion	Google
MT-NLG	530 Billion	Microsoft NVIDIA
PanGu- α	200 Billion	PengCheng
GPT-3	175 Billion	OpenAI
Turing-NLG	17.2 Billion	Microsoft

Treviso, Marcos, et al. "Efficient methods for natural language processing: A survey." *Transactions of the Association for Computational Linguistics* 11 (2023): 826-860.



LLM Full Finetune Predicament



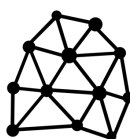
Source: Inside language models (from GPT-4 to PaLM) – Dr Alan D. Thompson – Life Architect



Large Language Model (LLM) enabled NLP applications

Pre-train

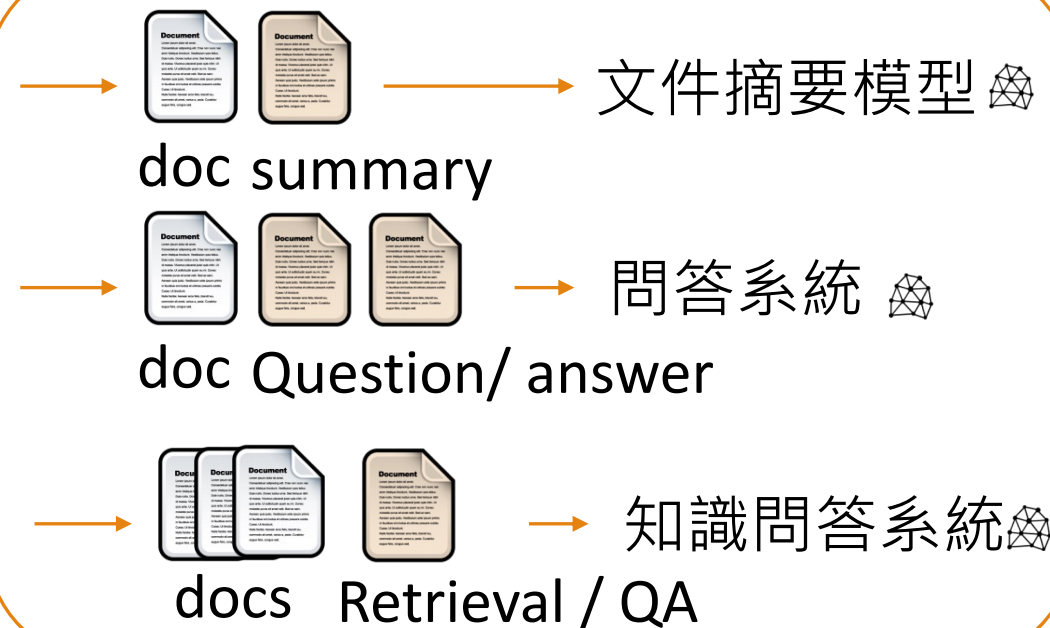




 $xB \sim xxxB$ 參數

BERT
GPT
ChatGPT
GPT-4

·
·
·

Fine-tune



GPU Memory Estimated – Full Finetune

Llama 2-7B, 16-bit float, seq 4096		
CUDA		~1Gb
Model weights	$\text{size(float)} * N_{\text{parameter}}$	13.03Gb
Gradients	$\text{size(float)} * N_{\text{trainable}}$	13.03Gb
Hidden states	$\sim \text{size(float)} L (20 \text{ seq} + 3 \text{ seq}^2)$	3.16Gb (batch size = 1)
Optimizer states	$2 * \text{size(float)} * N_{\text{trainable}}$	26.06Gb

L : Number of layers in model (eq. 32 layers)

H : Number of attention heads (eq. 32 heads)


Estimate: 56.28Gb



From Fine-tuning to Parameter-efficient Fine-tuning



Full Fine-tuning
Update **all model parameters**



Parameter-efficient Fine-tuning
Update a **small subset** of model parameters

GPU Memory Estimated – Train Less Parameters

Training only 0.2M parameters

Llama 2-7B, 16-bit float, seq 4096		
CUDA		~1Gb
Model weights	$\text{size(float)} * N_{\text{parameter}}$	13.03Gb
Gradients	$\text{size(float)} * N_{\text{trainable}}$	0.4Mb
Hidden states	$\sim \text{size(float)} L (20 H \text{ seq} + 3 \text{ seq}^2)$	3.16Gb (batch size = 1)
Optimizer states	$2 * \text{size(float)} * N_{\text{trainable}}$	0.8Mb

L : Number of layers in model (eq. 32 layers)

H : Number of attention heads (eq. 32 heads)

Estimate: 17.19Gb



Haven't We Seen This Before ?

- Updating the last layer was common in computer vision
- In NLP, people experimented with static and non-static word embeddings
- ELMo did not fine-tune contextualized word embeddings



Matthew E. Peters, et al. "Deep Contextualized Word Representations." Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers). 2018.




Benefits of PEFT

1. Decreased computational and storage costs

- One significant benefit of parameter-efficient fine-tuning lies in its reduced computational and storage demands compared to the high costs associated with full fine-tuning.

2. Portability


- It offers applicability across various domains and tasks. PEFT achieves this by **introducing a limited set of task-specific parameters while preserving the general-purpose parameters of the pre-trained model**, facilitating seamless transfer to new unlabeled datasets.



Benefits of PEFT

3. Overcoming catastrophic forgetting

- **When extensively pre-trained models undergo full fine-tuning on a novel task, it often leads to the model forgetting previously acquired knowledge** from its pre-training phase. The insights gained from the prior task are overridden by the updates tailored for the new task.
- Through PEFT, **only a small subset of parameters undergo modification during fine-tuning, leaving the majority of the model - which encapsulates general language knowledge from pre-training - unchanged.** This focused adjustment helps prevent the loss of knowledge from the original pre-training task.



Benefits of PEFT

4. Better performance in low-data regimes


- Large Language models of considerable size possess hundreds of millions or even tens of millions of trainable parameters, **thereby increasing the risk of overfitting when fine-tuned on tasks with limited labeled examples.**
- Nevertheless, **PEFT selectively trains only a small subset of these parameters, leveraging knowledge from the robust pre-trained model.** This approach enables the model to generalize more effectively as it still heavily relies on the extensive pre-trained representation.



Benefits of PEFT

5. Performance comparable to full fine-tuning

- **Extensive studies have demonstrated that parameter-efficient techniques can match or surpass the performance achieved by fully fine-tuning pre-trained models**, despite adjusting only a minute fraction of parameters.
- For instance, **research indicates that incorporating a small adapter module during fine-tuning yields performance results within 1% of fully adapting BERT across various natural language understanding benchmarks**. This illustrates that PEFT isn't just a workaround with compromised accuracy, but rather a method capable of achieving similarly robust outcomes while leveraging its numerous advantages over full fine-tuning.



Trainable Parameters Comparison

Method	RTE (Acc)	Trainable parameters (M)	Ratio
Full Fine-tune Acc	83.75%	184	100%
AdaLoRA	88.09%	1.27	0.69%
LoRA	86.60%	0.8	0.43%
Random Rank LoRA[†] (rank : 1 - 16)	85.56%	0.62	0.34%
Random Rank LoRA[†] (rank : 8 - 24)	85.16%	1.18	0.64%
Random Rank LoRA[†] (rank : 16 - 32)	85.13%	1.77	0.96%
Random Rank LoRA[†] (rank : 32 - 64)	84.91%	3.54	1.92%

[†] : Average of Samples (10 samples)

* Model: DeBERTa-v3-base

** Ratio means compare with Full Fine-tune




PEFT Theory

Intrinsic Dimensionality

The definition of Intrinsic Dimensionality (d_{int})

- **Intrinsic dimensionality refers to the dimensionality of the solution set within the overall parameter space**
 - If a high-dimensional space ($D = 1000$), the effective dimensionality ($d_a = 10$) means we optimize by random subspace ($d = 10$) can achieve “a %” performance of original optimizing outcome.




Li, Chunyuan, et al. "Measuring the Intrinsic Dimension of Objective Landscapes." *International Conference on Learning Representations*. 2018.



Intrinsic Dimensionality

Measuring the Intrinsic Dimension of Objective Landscapes

- Define d_{100} as the intrinsic dimension of the “100%” solution: solutions whose performance is statistically indistinguishable from baseline solutions



In this example, the d_{90} is approximately equal to 800.

Li, Chunyuan, et al. "Measuring the Intrinsic Dimension of Objective Landscapes." *International Conference on Learning Representations*. 2018.



Intrinsic Dimensionality

Many problems have smaller intrinsic dimensions than one might suspect

Dataset	MNIST		CIFAR-10		Inverted Pendulum	Humanoid	Atrai Pong
Network Type	FC	LeNet	FC	LeNet	FC	FC	ConvNet
Parameter Dim. D	199,210	44,426	656,810	62,006	562	166,673	1,005,974
Intrinsic Dim. d_{90}	750	290	9,000	2,900	4	700	6,000
d_{90} / D	0.38%	0.65%	1.37%	4.68%	0.7%	0.42%	0.60%


Li, Chunyuan, et al. "Measuring the Intrinsic Dimension of Objective Landscapes." *International Conference on Learning Representations*. 2018.



Intrinsic Dimensionality

Pre-training implicitly minimizes intrinsic dimension

- Compute d_{90} for six datasets: MRPC, QQP, Yelp Polarity, SST-2, MNLI, and ANLI
- See that **the intrinsic dimensionality of RoBERTa-base monotonically decreases as we continue pre-training**




Aghajanyan, Armen, Sonal Gupta, and Luke Zettlemoyer. "Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning." *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. 2021.



Intrinsic Dimensionality

Larger models tend to have lower intrinsic dimension after a fixed number of pre-training updates

- Used the MRPC dataset and computed intrinsic dimension for every pre-trained model
- See a strong general trend that **as the number of parameters increases, the intrinsic dimension of fine-tuning on MRPC decreases**



Aghajanyan, Armen, Sonal Gupta, and Luke Zettlemoyer. "Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning." *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. 2021.




Intrinsic Dimensionality

Observations

- Many problems have smaller intrinsic dimensions
- Intrinsic dimensionality decreases during pre-training
- Larger models have lower intrinsic dimensionality

- Parameters of model
- Pre-training updates



- Intrinsic dimension

Aghajanyan, Armen, Sonal Gupta, and Luke Zettlemoyer. "Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning." *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. 2021.



PEFT

Current Development & Method

The Venn diagram classifies LLM adaptation methods into four categories: additive (blue), selective (yellow), adapters (purple), and soft prompts (teal). The methods are distributed as follows:

- additive only:** Ladder-Side Tuning
- selective only:** BitFit, LN Tuning, Attention Tuning, Diff-Pruning, Fish-Mask, LT-SFT, FAR
- adapters only:** Parallel Adapters
- soft prompts only:** WARP, Spot, Prompt-tuning
- additive & selective:** Sparse Adapter
- additive & adapters:** Adapters, AdaMix
- additive & soft prompts:** LeTS, (IA)³, IPT, Prefix-Tuning
- selective & adapters:** S4
- selective & soft prompts:** KronA^{B_{res}}, PHM Adapter, Compacter
- adapters & soft prompts:** MAM Adapter, UniPELT
- additive, selective, & adapters:** AttentionFusion
- additive, selective, & soft prompts:** LeTS
- additive, adapters, & soft prompts:** AdaMix
- selective, adapters, & soft prompts:** S4
- additive, selective, & adapters, & soft prompts:** AdaMix
- additive & selective & adapters:** Parallel Adapters
- additive & selective & soft prompts:** LeTS
- additive & adapters & soft prompts:** AdaMix
- selective & adapters & soft prompts:** S4
- additive, selective, adapters, & soft prompts:** AdaMix


Parameter Efficient Fine-Tuning

PEFT Current Development

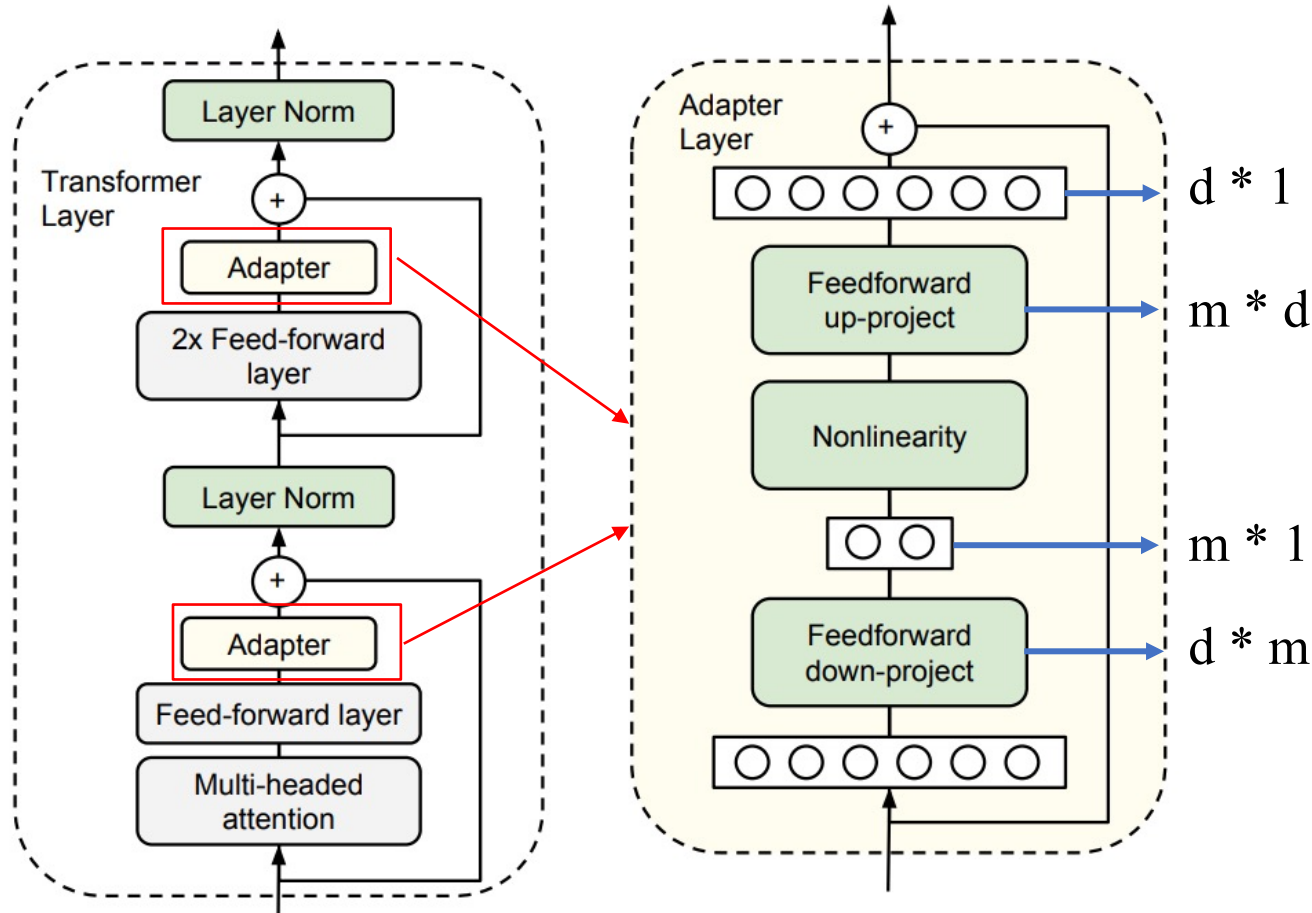
- **Additive** Method:
 - Additive fine-tuning approaches involve introducing new extra trainable parameters for task-specific fine-tuning.
- **Selective** Method:
 - Selective fine-tuning methods aim to reduce the number of fine-tuned parameters by selecting a subset of pre-trained parameters that are critical to downstream tasks while discarding unimportant ones.
- **Reparametrization-based** Method:
 - Reparameterized fine-tuning methods utilize low-rank transformation to reduce the number of trainable parameters while allowing operating with high-dimensional matrices (e.g., pretrained weights).



PEFT Current Development



Additive PEFT: Adapters



- The adapters first project the original d -dimensional features into a smaller dimension, m , apply a nonlinearity, then project back to d dimensions.
- By setting $m \ll d$, we limit the number of parameters added per task

Houlsby, Neil, et al. "Parameter-efficient transfer learning for NLP." *International Conference on Machine Learning*. PMLR, 2019.

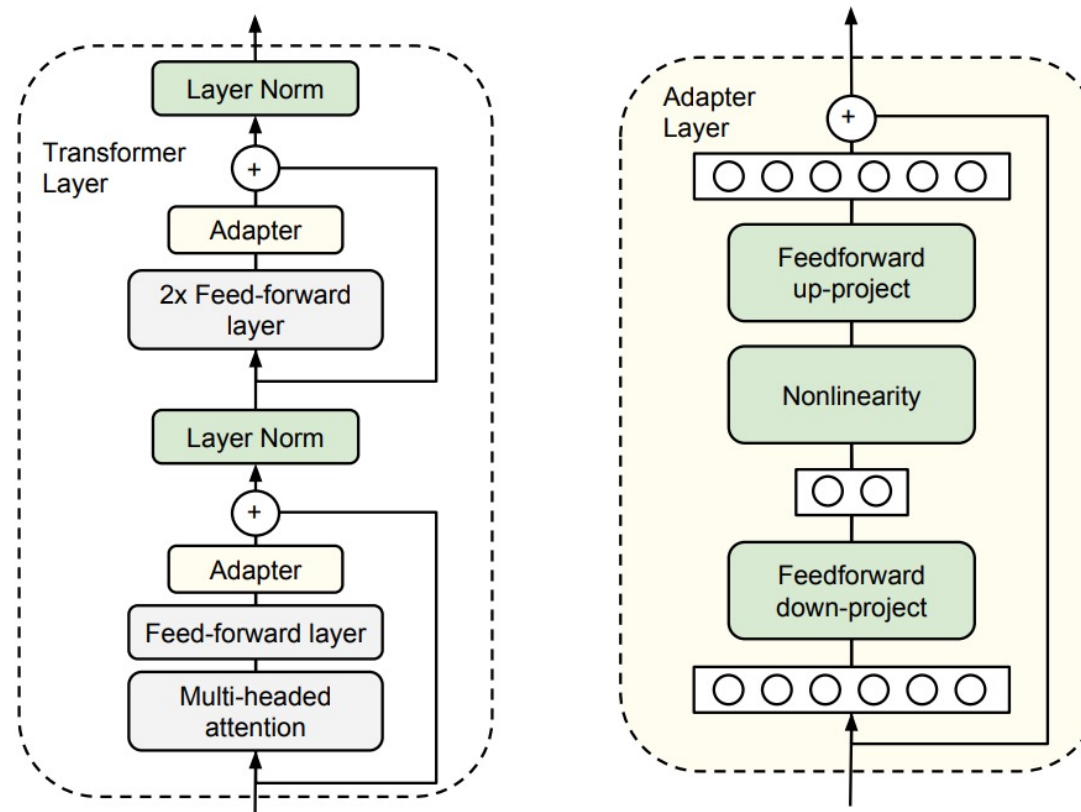


Additive PEFT: Adapters

Add fully-connected networks after attention and FFN layers

Pseudocode:

```
def transformer_block_with_adapter(x):  
    residual = x  
    x = SelfAttention(x)  
    x = FFN(x) # adapter  
    x = LN(x + residual)  
    residual = x  
    x = FFN(x) # transformer FFN  
    x = FFN(x) # adapter  
    x = LN(x + residual)  
    return x
```

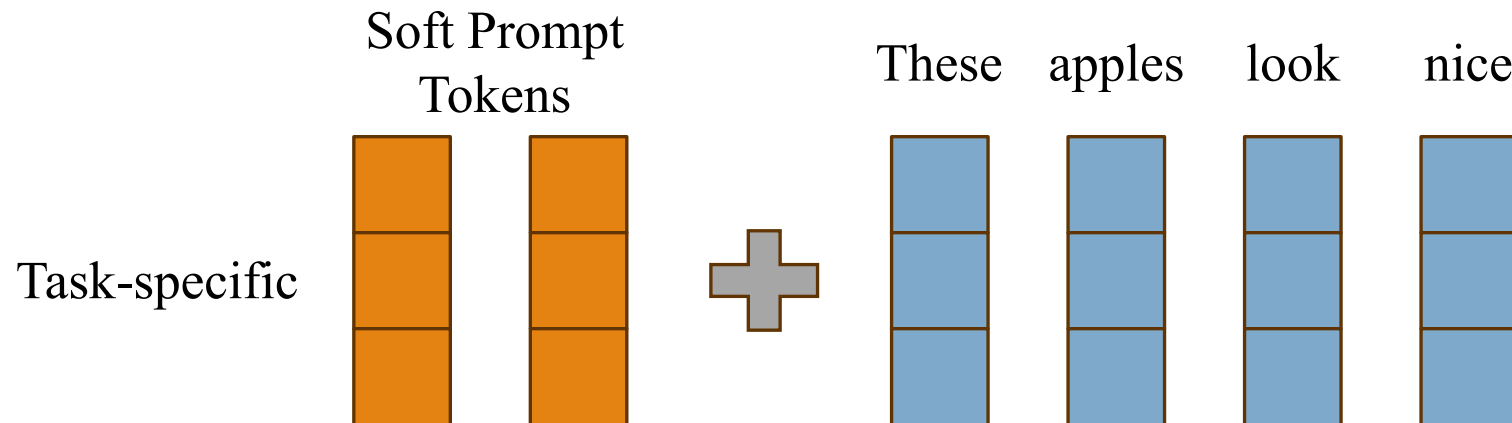


Houlsby, Neil, et al. "Parameter-efficient transfer learning for NLP." *International Conference on Machine Learning*. PMLR, 2019.



Additive PEFT: Prompt Tuning

- Concatenates trainable parameters with the input embeddings
 - Learn a new sequence of task-specific embeddings
 - We call this prompt tuning because we only update prompt weights



Lester, Brian, Rami Al-Rfou, and Noah Constant. "The Power of Scale for Parameter-Efficient Prompt Tuning." *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.



Additive PEFT: Prompt Tuning

- Efficient for multi-task serving
 - Each task is a prompt, not a model
 - Prompts for various tasks can be applied to different inputs

	Soft Prompt Tokens		Input embedding vectors			
Sentiment	[.6,...,-4.3]	[.2,...,5.4]	These	apples	look	nice
Q&A	[-1.2,...,8]	[1.3,...,-2.7]	When	should	I	leave
Translate	[-.5,...,-1.3]	[.9,...,-.5]	I	love	fresh	air

Lester, Brian, Rami Al-Rfou, and Noah Constant. "The Power of Scale for Parameter-Efficient Prompt Tuning." *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.



Additive PEFT: Prompt Tuning

Prepend the model input embeddings with a trainable tensor


Pseudocode:

```
# make it a parameter so that it can be trained
# Initialize soft_prompt tensor with random values
soft_prompt = torch.nn.parameter(
    torch.rand(num_tokens, embedding_dim)
)
# Concatenate soft_prompt with input x
def input_soft_prompt(x, soft_prompt):
    x = concatenate([soft_prompt, x],
                    dim = seq_len)

    return x

# train soft_prompt tensor with gradient descent
train(model(input_soft_prompt(x, soft_prompt)))

# use model with soft_prompt
model(input_soft_prompt(x, soft_prompt))
```



Lester, Brian, Rami Al-Rfou, and Noah Constant. "The Power of Scale for Parameter-Efficient Prompt Tuning." *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.

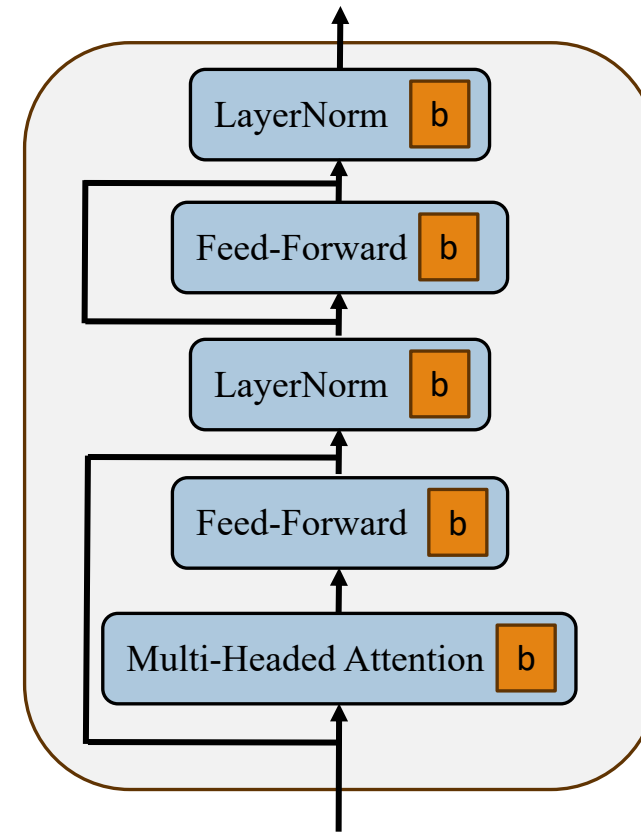


Selective PEFT: BitFit

Fine-tune only model biases

Pseudocode:

```
params = (p for n,p
          in model.named_parameters()
          if "bias" in n)
optimizer = Optimizer(params)
```



Zaken, Elad Ben, Yoav Goldberg, and Shauli Ravfogel. "BitFit: Simple Parameter-efficient Fine-tuning for Transformer-based Masked Language-models." *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. 2022.




Reparametrization-based PEFT: LoRA

$$\begin{aligned} h &= W_0 x + \alpha \Delta W x \\ &= W_0 x + \alpha B A x \end{aligned}$$

$$A \in \mathbb{R}^{r \times k} \quad B \in \mathbb{R}^{d \times r}$$

Matrix rank Hidden dimension Input dimension

$$\alpha \in \mathbb{R}^+ \leftarrow \text{Constant}$$



Hu, Edward J., et al. "LoRA: Low-Rank Adaptation of Large Language Models." *International Conference on Learning Representations*. 2021.



Reparametrization-based PEFT: LoRA

Decompose a weight matrix into lower-rank matrices

Pseudocode:

```
input_dim = 768 # the hidden size of the pre-trained model
output_dim = 768 # the output size of the layer
rank = 8 # The rank 'r' for the low-rank adaptation


W = ... # from pretrained network with shape input_dim x
output_dim

W_A = nn.Parameter(torch.empty(input_dim, rank)) # LoRA weight A
W_B = nn.Parameter(torch.empty(rank, output_dim)) # LoRA weight B

# Initialization of LoRA weights
nn.init.kaiming_uniform_(W_A, a=math.sqrt(5))
nn.init.zeros_(W_B)

def regular_forward_matmul(x, W):
    h = x @ W
    return h

def lora_forward_matmul(x, W, W_A, W_B):
    h = x @ W # regular matrix multiplication
    h += x @ (W_A @ W_B) * alpha # use scaled LoRA weights
    return h
```

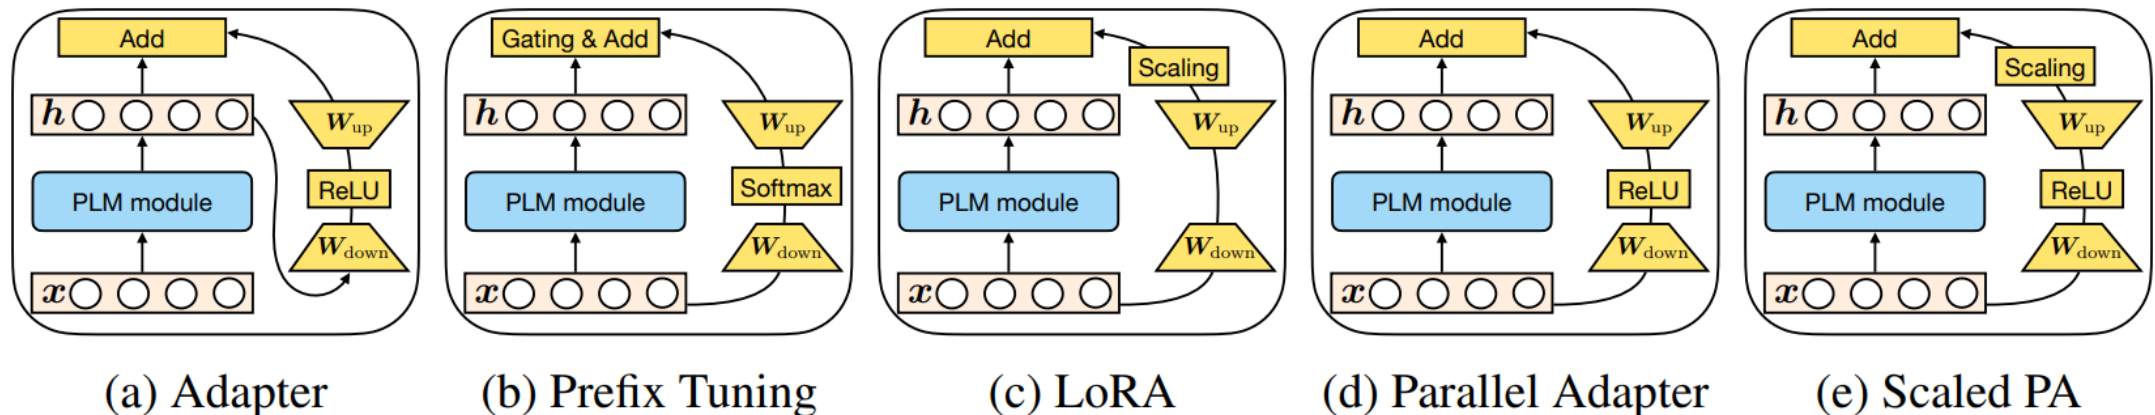


Hu, Edward J., et al. "LoRA: Low-Rank Adaptation of Large Language Models." *International Conference on Learning Representations*. 2021.



Hybrid PEFT: MAM Adapters

- Break down the design of state-of-the-art parameter-efficient transfer learning methods and present a unified framework that establishes connections between them.



He, Junxian, et al. "Towards a Unified View of Parameter-Efficient Transfer Learning." *International Conference on Learning Representations*. 2021.




Hybrid PEFT: MAM Adapters

Scaled parallel adapter for FFN layer + soft prompt

Pseudocode:

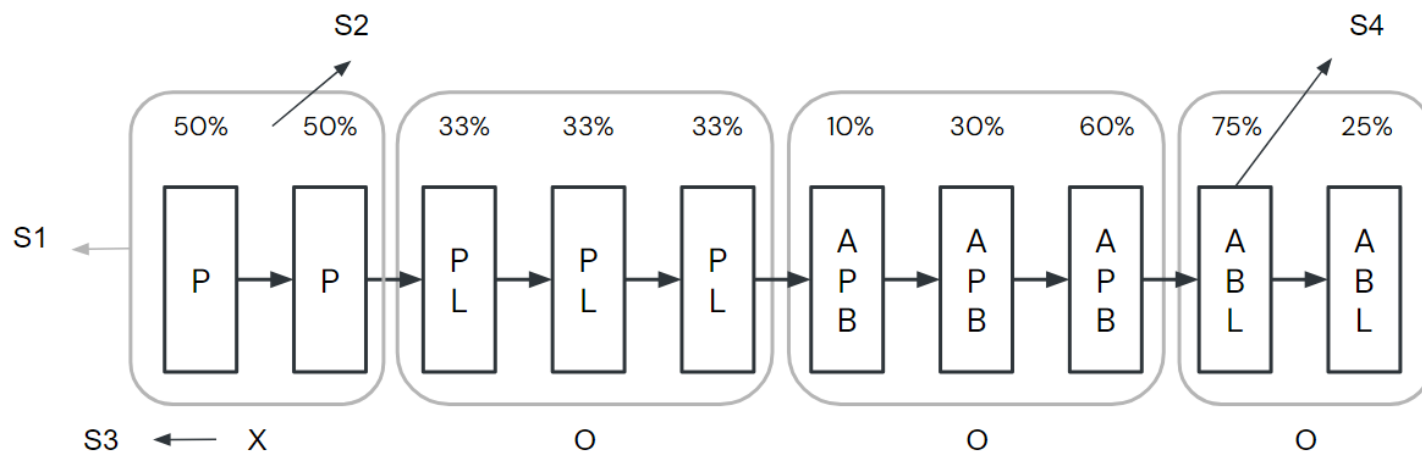
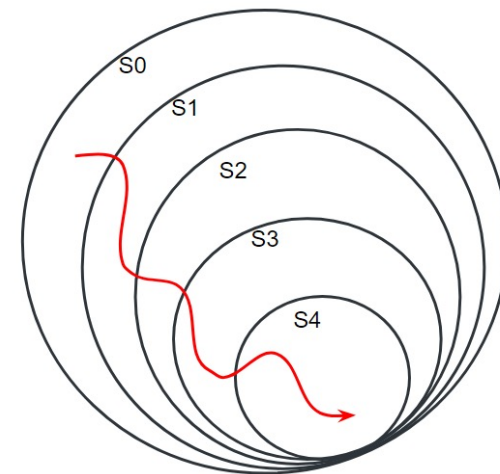
```
def transformer_block_mam(x):  
    x = concat([x, soft_prompt],  
               dim=seq)  
  
    residual = x  
    x = SelfAttention(x)  
    x = LN(x + residual)  
    x_a = FFN(x) # parallel adapter  
    x_a = scale * x_a  
    x = LN(x + x_adapter)  
    return x
```



He, Junxian, et al. "Towards a Unified View of Parameter-Efficient Transfer Learning." *International Conference on Learning Representations*. 2021.

Hybrid PEFT: S4

- Designing Network Design Spaces
 - S0_The Initial Design Space: a random strategy
 - S1_Layer Grouping: Increasing, Uniform, Decreasing, Spindle, Bottleneck
 - S2_Trainable Parameter Allocation: Increasing, Uniform, Decreasing
 - S3_Tunable Groups: Tune or not
 - S4_Strategy Assignment: {Adapter (A), Prefix (P), BitFit (B), and LoRA (L)}



Chen, Jiaao, et al. "Parameter-Efficient Fine-Tuning Design Spaces." *The Eleventh International Conference on Learning Representations*. 2022.


Hybrid PEFT: S4

Automatically found combination of Adapters, Prefix-Tuning, BitFit, and LoRA

Search process:

At 0.1% additional parameters

1. Find optimal grouping pattern: spindle
2. Find optimal parameter allocation pattern: uniform
3. Find which groups need tuning: all
4. find optimal combinations of PEFT techniques sequentially for G1, G2, G3, G4



$$G_1 : A, L \quad G_3 : A, P, B$$

$$G_2 : A, P \quad G_4 : P, B, L$$

PEFT Method: Summary

Method	Type	Storage	Memory	Inference overhead	Trainable parameters	Changed parameters
Adapters	A	yes	yes	Extra FFN	0.1% - 6%	0.1% - 6%
Prompt Tuning	A	yes	yes	Extra input	0.1%	0.1%
Bitfit	S	yes	yes	Extra input	0.5%	0.5%
LoRA	R	yes	yes	No overhead	0.01% - 0.5%	0.5% - 30%
MAM Adapters	A	yes	yes	Extra FFN & input	0.5%	0.5%
S4	ASR	yes	yes	Extra FFN & input	0.5%	> 0.5%

Lialin, Vladislav, Vijeta Deshpande, and Anna Rumshisky. "Scaling down to scale up: A guide to parameter-efficient fine-tuning." *arXiv preprint arXiv:2303.15647* (2023).



What Matters in PEFT ?

1. Number of parameters
2. Training efficiency
 - Do we add parameters to the network? How expensive are they?
 - Does the method require to backpropagate through the original network?
 - Can the method efficiently utilize the GPU?
3. Inference efficiency
 - Do we add parameters to the network? How expensive are they?
4. Accuracy
 - Do we get good performance out of the network in the end?



Appendix

Formula for Hidden States Estimate

During training:

$$(3 \text{ h seq} + L(4 \text{ h seq} + 3 \text{ h seq}^2 + 8 \text{ h seq} + 2 \text{ h seq} + 4 \text{ h seq}) + \text{vocab seq}) \text{ bs} =$$

Emb,pos, Pre-logit h	K,V,Q,O	Score, Probs, Dropout	FFN hidden, activation	FFN out, dropout	residual, LN	logits
-------------------------	---------	-----------------------------	---------------------------	---------------------	-----------------	--------

$$= 3 \text{ h seq bs} + 18L \text{ h seq bs} + 3L \text{ heads seq}^2 + \text{vocab seq bs}$$

L : Number of layers in model (eq. 32 layers)

H : Number of attention heads (eq. 32 heads)

bs : batch size

